


Introduction

Proteins are the molecular workhorses of life, orchestrating virtually every biological process from catalyzing reactions to transmitting signals and maintaining cellular structure. Understanding protein function from amino acid sequence alone remains one of the most fundamental challenges in biology. This challenge has direct implications across medicine, industry, and basic research, yet the vast majority of known proteins remain functionally uncharacterized.


Applications

Understanding protein function has profound implications across virtually every domain of biological research and biotechnology.



The Dark Proteome

Despite the critical importance of understanding protein function, a massive gap exists between available sequence data and functional knowledge—creating a fundamental bottleneck limiting human progress against disease, environmental challenges, and industrial applications.



Traditional Computational Approaches

Consequently, the field has turned to computational prediction as the only scalable solution to bridge this widening gap between sequence availability and functional understanding.

BLAST: For over three decades, the most widely used tool for protein annotation; identifies sequences similar to already-characterized proteins and transfers their functional labels via homology-based inference

Limitations of homology transfer:

- Functional annotation is typically shallow even when similarity is plausible
- Can only annotate proteins that closely resemble already-characterized sequences
- Leaves the vast majority of the dark proteome—especially proteins from understudied organisms, extreme environments, and metagenomic discoveries—functionally invisible

Current Deep Learning Methods

More recently, deep learning methods have shown promise by learning latent relationships between sequences and functions that transcend simple sequence similarity, yet these approaches face their own critical limitations. Three key limitations persist.

Underutilization of protein language models (PLMs)

- Prevailing systems compress rich embeddings through aggressive low-dimensional bottlenecks, rely on a single PLM despite complementary inductive biases across architectures, or employ prediction frameworks that fail to fully express PLM representational capacity

Misattribution of performance drivers

- The field frequently overemphasizes structural inputs and places undue weight on auxiliary evolutionary information such as multiple sequence alignments (MSAs)

Auxiliary feature dependencies

- Requirements for predicted structures from computationally expensive pipelines (e.g., AlphaFold) or evolutionary alignments from database searches add substantial computational overhead, increase inference latency, limit applicability to novel proteins lacking evolutionary history, and risk propagating errors from upstream predictions

Research Objectives

Investigate the true drivers of performance gains in protein function prediction through systematic ablation experiments—isolating the contributions of auxiliary inputs (3D coordinates, MSAs) versus protein language model representations

Design a PLM-centric architecture that fully leverages complementary protein language models through advanced fusion techniques (cross-attention, adaptive gating) and training strategies validated through systematic experimentation


Evaluate rigorously on Gene Ontology prediction across all three ontologies (MF, BP, CC) against state-of-the-art methods, with additional benchmarking on enzyme classification and stability regression to assess generalization and fine-tuning capability

Develop an accessible, serverless web application with a user-friendly interface for sequence-to-function prediction, alongside open-source implementation enabling mass annotation at scale

Protein Data Modality

Proteins are linear polymers defined entirely by their amino acid sequence—a chain of residues drawn from a 20-letter alphabet that encodes all the information necessary to determine a protein's structure, function, and behavior. This primary sequence is the most fundamental and universally available representation of any protein.

VRDYAIKPCNCVEYCGITDCKNTLCTENGAESGYCQWGGYKGNACWCICKLPDSVPVIRPGKCKQ



Datasets & Evaluation Metrics

Gene Ontology

Task Background

- The Gene Ontology (GO) provides a controlled vocabulary for describing protein function across three orthogonal ontologies: molecular function (MF), biological process (BP), and cellular component (CC)
- MF captures specific biochemical activities (e.g., ATP binding, oxidoreductase activity); BP describes broader pathways (e.g., DNA repair, signal transduction); CC localizes proteins to subcellular structures (e.g., cytosol, ribosome)
- Technically challenging: labels are hierarchical, exhibit long-tailed frequency distributions, and require models to generalize to remote homologs and rare terms while remaining ontology-consistent
- Central task targeted in this work—regarded as one of the most important benchmarks in protein informatics since it is well-defined while addressing the key overall aspects of protein function through its three ontologies

Dataset

- PDBCh: 36,641 protein chains from the Protein Data Bank clustered at 95% sequence identity, split 8:1:1 into 29,313 training, 3,664 validation, and 3,664 test proteins
- Label space: 489 MF terms, 1,943 BP terms, 320 CC terms—exhibiting moderate label imbalance typical of real-world annotation databases (rare terms ~50 examples, common terms in thousands)
- Test set remains fixed across all three ontologies for standard cross-study comparability

Evaluation Metrics

- GO predictions undergo standard hierarchy-aware propagation preprocessing to ensure consistency with the GO ontology tree
- Protein-centric F_{max}** : Best thresholded trade-off between per-protein precision and recall; higher means balanced, reliable hits at the protein level
- Function-centric $AUPR_{macro}$** : Averages average-precision over GO terms so rare and common labels count equally; higher means better performance on long-tailed terms
- Semantic distance S_{min}** : Minimizes information-weighted error on the GO DAG; lower values indicate more semantically accurate and hierarchy-consistent predictions

$$F_{max} = \max_{\tau \in \mathcal{T}} \frac{2\bar{P}(\tau)\bar{R}(\tau)}{\bar{P}(\tau) + \bar{R}(\tau)} \quad AUPR_{macro} = \frac{1}{|\mathcal{L}^+|} \sum_{t \in \mathcal{L}^+} AP_t \quad S_{min} = \min \sqrt{RU(\tau)^2 + MI(\tau)^2}$$

Enzyme Commission

Task Background

- EC number: hierarchical four-level taxonomy specifying the chemical transformation catalyzed by an enzyme (class $\rightarrow$ subclass $\rightarrow$ sub-subclass $\rightarrow$ substrate specificity)
- Near single-label classification with extremely long-tailed frequency distribution

Dataset

- Training: 74,487 sequences clustered at 70% identity from UniProt; Validation: NEW-392 (392 Swiss-Prot enzymes); Test: Price-149 (149 experimentally validated enzymes)

Evaluation Metric

- AUC: Area under ROC curve measuring class separability; higher means stronger discrimination

$$AUC = \int_0^1 TPR(FPR) d(FPR), \quad TPR = \frac{TP}{TP + FN}, \quad FPR = \frac{FP}{FP + TN}$$

Stability

Task Background

- Protein stability: ability of a sequence to retain its native fold under stress (heat, denaturation, proteolysis), summarized as a continuous score
- Regression task probing mutational landscape; models must generalize beyond single-mutation trends

Dataset

- TAFE stability benchmark: 53,416 training, 2,512 validation, 12,851 test sequences from a high-throughput de novo miniprotein design campaign

Evaluation Metric

- Spearman's ρ : Rank correlation between predicted and true stability scores; higher means stronger monotonic agreement

$$\rho = \frac{\text{Cor}(r(s), r(y))}{\sigma_{r(s)}\sigma_{r(y)}} = 1 - \frac{6 \sum_{i=1}^n (r(s_i) - r(y_i))^2}{n(n^2 - 1)}$$

Input Modality Investigation

Multiple Sequence Alignments

Background & Motivation

- Multiple sequence alignments (MSAs) capture evolutionary relationships by aligning homologous sequences, revealing conservation patterns and co-evolutionary signals
- Key question: Do MSAs provide measurable benefit to protein function prediction tasks when strong protein language models are available?

Experimental Setup

- Compared two architecturally identical models on MF prediction with varied inputs:
- Baseline**: Four PLM embeddings (ESM-C, ProtT5, Ankh3-XL, PGLM)
- MSA-augmented**: Pooled MSA features replace PGLM as fourth stream
- Substitution design ensures identical capacity, isolating effect of auxiliary evolutionary information vs single sequence modeling
- MSAs encoded via MSA Transformer with learned conservation-weighted pooling (see diagram):

$$x_{rel} = W^T M_{rel}, \quad \sigma_{rel} = \frac{\exp(s_{rel}(\tau\sqrt{N}))}{\sum_{s' \in \mathcal{S}} \exp(s'(\tau\sqrt{N}))}, \quad h_t = \sum_{s \in \mathcal{S}} \sigma_{s,t} M_{s,t}$$

Data Generation

- Retrieved pre-computed MSAs from OpenFold dataset (BFD+UniClust30, MGnify, UniRef90) with Hamming diversity-maximizing subsampling
- For proteins absent from OpenFold: generated MSAs using ColabFold's MMSeqs2 protocol (90% identity threshold, 75% coverage, max 256 sequences)

Results & Interpretation

- Final validation $AUPR_{micro}$: 0.877 (without MSA) vs 0.872 (with MSA); MSA features provide minimal measurable benefit when protein language models exist as a prior
- Interpretation: Modern PLMs pretrained on millions/billions of sequences have internalized relevant co-evolutionary patterns necessary for protein function prediction, making explicit MSA input redundant

Structural Information

A. KNN GRAPH CONSTRUCTION

Real Coordinates vs Random Coordinates

B. EGGN LAYER (*4)

Distance Features: $d_{ij} = \|c_i - c_j\|^2$

C. READOUT PIPELINE

Processed Node Features (N) $\rightarrow$ Masked Mean Pooling $g = \sum_i h_i / N$ $\rightarrow$ Protein Embedding $g \otimes P$ $\rightarrow$ Classification Head (MF/BP/CC predictions)

D. Ablation Results

Scaling Protein Language Models

Background & Motivation

- Nearly all existing prediction methods rely on a single PLM to extract sequence features
- Key question: Does incorporating multiple PLMs with diverse architectural biases improve performance?

Epoch	1	2	3	4	6
1	0.146	0.272	0.287	0.341	0.355
2	0.389	0.540	0.601	0.646	0.644
3	0.553	0.675	0.744	0.760	0.777
4	0.632	0.745	0.804	0.809	0.822
5	0.710	0.795	0.842	0.845	0.845
6	0.754	0.815	0.856	0.861	0.861
7	0.791	0.837	0.867	0.877	0.875


Experimental Setup

- Systematically varied number of PLM streams on MF prediction over first 7 epochs
- Starting from ESM-C baseline, incrementally added: ProtT5 $\rightarrow$ Ankh3-XL $\rightarrow$ PGLM $\rightarrow$ MSA Transformer $\rightarrow$ GENIE
- Architecture adapts to PLM count: 1 PLM = self-attention encoder; 2 PLMs = bidirectional cross-attention; 3+ PLMs = dual-stream encoder with grouped representations
- All variants hold fixed modules to control for model capacity

Results & Interpretation

- Performance improves monotonically from 1 $\rightarrow$ 4 PLMs (epoch 7: 0.791 $\rightarrow$ 0.837 $\rightarrow$ 0.867 $\rightarrow$ 0.877)
- Beyond 4 PLMs, gains plateau (6 PLMs: 0.875)—diminishing returns from redundancy
- Interpretation: Each PLM contributes distinct inductive biases from its architecture, pretraining corpus, and training objective—complementary representations drive performance gains
- Optimal sweet spot: 3-4 complementary PLMs for protein representation learning

Architecture



Overview

- Four complementary PLMs with distinct inductive biases: ESM-C (efficiency-scaled MLM), ProtT5 (encoder-decoder long-range dependencies), Ankh3-XL (multi-task pretraining), PGLM (unified autoencoding + autoregressive)
- Dual-Stream Encoder: PLMs grouped into two streams (ESM-C + Ankh3 | ProtT5 + PGLM)—intra-stream fusion via bidirectional cross-attention with convex gating produces per-stream representations
- Bidirectional Cross-Attention Transformer: Streams exchange information through N layers of symmetric cross-attention (each stream attends to the other) followed by SwiGLU feed-forward networks
- Masked Mean Pooling: Aggregates token-level representations to protein-level embeddings, excluding padding and BOS/EOS registers
- Convex Gated Fusion: Learned softmax-normalized weights (summing to 1) combine stream embeddings—ensuring stable mixing
- MLP Prediction Head: Task-specific output layer produces GO terms, EC numbers, stability scores, or any other relevant objective
- Key design principle: enable PLMs to communicate and share complementary insights via cross-attention—allowing selective querying of relevant features across streams—rather than relying on simple concatenation, while incorporating cutting-edge mathematical techniques and disregarding redundant auxiliary inputs

Training Techniques

- Loss Function**
 - Asymmetric Loss (ASL) for multi-label classification (GO, EC)—focal-style objective addressing extreme class imbalance:
$$\mathcal{L}_{ASL}(\mathbf{x}, \mathbf{y}) = -\sum_{c=1}^C \{y_c \cdot \log(p_c^+) + (1 - y_c) \cdot \log(p_c^-)\} \cdot w_c$$
$$w_c = (1 - p_c^+)^{\gamma_c} \cdot y_c - p_c^- \cdot (1 - y_c)^{\gamma_c}, \quad \gamma_c = \gamma_{pos} \cdot y_c + \gamma_{neg} \cdot (1 - y_c)$$
 - Standard Mean Squared Error (MSE) loss for stability regression:
$$\mathcal{L}_{MSE}(\mathbf{x}, \mathbf{y}) = \frac{1}{N} \sum_{i=1}^N (x_i - y_i)^2$$
- Optimization**
 - AdamW for most efficient convergence and loss exploration
 - Learning rate $\eta = 10^{-4}$ for parameters, weight decay $\lambda = 0.01$
 - Momentum: $\beta_1 = 0.9$, $\beta_2 = 0.999$, numerical stability $\epsilon = 10^{-8}$
 - Gradient clipping at norm 1.0 to prevent training instability
- Learning Rate Scheduler**
 - Cosine annealing with linear warmup (5% of training steps)
 - Warmup stabilizes early optimization; cosine decay ensures smooth convergence
 - $\eta_t = \begin{cases} \eta_{warmup} & \text{if } t < T_{warmup} \text{ (linear warmup)} \\ \eta \cdot \frac{1}{2} \left(1 + \cos\left(\pi \cdot \frac{t - T_{warmup}}{T_{train} - T_{warmup}}\right)\right) & \text{otherwise (cosine decay)} \end{cases}$
- Weights Averaging Scheme**
 - Exponential Moving Average (decay $\tau = 0.999$), improves model robustness and generalization
 - EMA accumulation begins after warmup; EMA weights used at inference
 - $\theta_{EMA}^{(t)} = \tau \cdot \theta^{(t-1)} + (1 - \tau) \cdot \theta^{(t)}$
- Early Stopping**
 - Monitors validation performance with patience of 2 epochs
 - Minimum improvement threshold $\Delta_{min} = 0.003$
 - Tracks $AUPR_{macro}$ for classification tasks, Spearman's ρ in regression

Ablations

Epoch	1	2	3	4	5	6	7
Bidirectional Transformer vs Self-Attention Transformer							
Bidirectional Transformer	0.3372	0.6382	0.7592	0.8180	0.8358	0.8638	0.8724
Self-Attention Transformer	0.3318	0.6180	0.7288	0.7895	0.8129	0.8383	0.8589
$\Delta(A-B)$	+0.0054	+0.0202	+0.0304	+0.0285	+0.0229	+0.0255	+0.0135
Cross Attention Pooling vs Naive Concatenation Pooling							
Cross Attention Pooling	0.2695	0.5756	0.7309	0.7841	0.8282	0.8490	0.8649
Naive Concatenation Pooling	0.2518	0.5364	0.6960	0.7764	0.8167	0.8305	0.8521
$\Delta(A-B)$	+0.0177	+0.0392	+0.0340	+0.0078	+0.0114	+0.0185	+0.0128
SwiGLU + Register Tokens							
With SwiGLU + Registers	0.2866	0.6012	0.7439	0.8036	0.8418	0.8556	0.8669
No SwiGLU/Registers	0.2518	0.5364	0.6960	0.7764	0.8167	0.8305	0.8521
$\Delta(A-B)$	+0.0348	+0.0649	+0.0479	+0.0273	+0.0250	+0.0251	+0.0148
Convex Gating vs Per-Dim Gating							
Convex Gating	0.2866	0.6012	0.7439	0.8036	0.8418	0.8556	0.8669
Per-Dim Gating	0.3143	0.6316	0.7581	0.8186	0.8358	0.8586	0.8592
$\Delta(A-B)$	-0.0277	-0.0304	-0.0141	-0.0160	-0.0060	-0.0030	-0.0076
High-Dimensional Embedding vs Low-Dimensional Embedding							
High-Dimensional Embedding	0.3408	0.6464	0.7596	0.8088	0.8453	0.8609	0.8770
Low-Dimensional Embedding	0.2995	0.6229	0.7459	0.8069	0.8425	0.8599	0.8724
$\Delta(A-B)$	+0.0413	+0.0235	+0.0137	+0.0019	+0.0028	+0.0010	+0.0045
Experiment Setup							
Systematic ablations on MF prediction task to validate architectural design choices							
All models trained for 7 epochs and evaluated using validation $AUPR_{micro}$							
Five architectural dimensions tested:							
Attention mechanism: Bidirectional cross-attention between streams vs independent self-attention within each stream							
Pooling strategy: Cross-attention streamed pooling vs naive concatenated projection pooling							
Architectural enhancements: SwiGLU activations + learnable register tokens vs standard configurations							
Gating strategy: Convex (softmax-normalized) vs per-dimension element-wise gating							
Embedding dimensionality: High-dimensional (Ankh3-XL, 25600) vs low-dimensional (Ankh3-L, 15360)							
Collectively: Design choices are complementary and compound to form a well-calibrated architecture							
Additional experiments (omitted for brevity): loss formulation (ASL outperformed BCE for handling extreme class imbalance), optimal latent dimensionality (896 balanced capacity and regularization), weight-averaging strategy (EMA improved inference smoothness over standard weights), learning rate schedules (cosine annealing with warmup stabilized convergence), optimizer comparisons, etc.							

Algorithm

Algorithm 1 Omni-Pro: PLM-Centric Fusion for Protein Function Prediction
Require: Amino acid sequence $s = (s_1, \dots, s_L)$
Ensure: Prediction vector $\mathbf{y} \in \mathbb{R}^C$

Stage 1: PLM Feature Extraction

- for $m \in \{\text{ESM-C, ProtT5, Ankh3-XL, PGLM}\}$ do
- $\mathbf{H}^{(m)} \leftarrow \text{Embed}(s)$
- end for

Stage 2: Projection to Shared Latent Space

- for each PLM embedding $\mathbf{H}^{(m)}$ do
- $\mathbf{Z}^{(m)} \leftarrow \text{Dropout}(\text{Proj}(\text{LN}(\mathbf{H}^{(m)}))) \in \mathbb{R}^{L \times d}$
- Append $\{\text{BOS}\}/\{\text{EOS}\}$ register tokens to $\mathbf{Z}^{(m)}$
- end for

Stage 3: Dual-Stream Encoder (Intra-Stream Fusion)

- for $(i, j) \in \{\{\text{ESM-C, Ankh3-XL}\}, \{\text{ProtT5, PGLM}\}\}$ do
- $\mathbf{Z}^{(i)} \leftarrow \text{MHA}(\text{LN}(\mathbf{Z}^{(i)}), \text{LN}(\mathbf{Z}^{(j)}), \text{LN}(\mathbf{Z}^{(i)}))$
- $\mathbf{Z}^{(j)} \leftarrow \text{MHA}(\text{LN}(\mathbf{Z}^{(j)}), \text{LN}(\mathbf{Z}^{(i)}), \text{LN}(\mathbf{Z}^{(j)}))$
- $\mathbf{W}^{(i,j)} \leftarrow \text{softmax}(\text{MLP}_{\text{gate}}(\mathbf{Z}^{(i)} \oplus \mathbf{Z}^{(j)}))$
- $\mathbf{Z}^{(i)} \leftarrow \mathbf{W}^{(i,j)} \odot \mathbf{Z}^{(i)} + \mathbf{W}^{(j,i)} \odot \mathbf{Z}^{(j)}$
- end for
- $\mathbf{X}^{(i)} \leftarrow \text{FF}_{\text{ESM-C, Ankh3-XL}}(\mathbf{Z}^{(i)})$
- $\mathbf{Y}^{(j)} \leftarrow \text{FF}_{\text{ProtT5, PGLM}}(\mathbf{Z}^{(j)})$

Stage 4: Bidirectional Cross-Attention Transformer

- for $\ell = 1, 2$ do
- $\mathbf{X}^{(\ell)} \leftarrow \mathbf{X}^{(\ell-1)} + \text{MHA}(\text{LN}(\mathbf{X}^{(\ell-1)}), \text{LN}(\mathbf{Y}^{(\ell-1)}), \text{LN}(\mathbf{Y}^{(\ell-1)}))$
- $\mathbf{Y}^{(\ell)} \leftarrow \mathbf{Y}^{(\ell-1)} + \text{MHA}(\text{LN}(\mathbf{Y}^{(\ell-1)}), \text{LN}(\mathbf{X}^{(\ell-1)}), \text{LN}(\mathbf{X}^{(\ell-1)}))$
- $\mathbf{X}^{(\ell)} \leftarrow \mathbf{X}^{(\ell)} + \text{FF}_{\text{SwiGLU}}(\text{LN}(\mathbf{X}^{(\ell-1)}))$
- $\mathbf{Y}^{(\ell)} \leftarrow \mathbf{Y}^{(\ell)} + \text{FF}_{\text{SwiGLU}}(\text{LN}(\mathbf{Y}^{(\ell-1)}))$
- end for

Stage 5: Masked Mean Pooling

- $\mathbf{X}_{\text{pool}} \leftarrow \sum_{i=1}^L \frac{M_i \mathbf{X}_i^{(N)}}{\sum_{i=1}^L M_i} + \mathbf{c}$
- $\mathbf{Y}_{\text{pool}} \leftarrow \sum_{j=1}^L \frac{M_j \mathbf{Y}_j^{(N)}}{\sum_{j=1}^L M_j} + \mathbf{c}$

Stage 6: Convex Gated Fusion

- $[\alpha, \beta] \leftarrow \text{softmax}(\text{MLP}_{\text{fusion}}(\mathbf{X}_{\text{pool}} \oplus \mathbf{Y}_{\text{pool}}))$
- $\mathbf{z} \leftarrow \alpha \cdot \mathbf{X}_{\text{pool}} + \beta \cdot \mathbf{Y}_{\text{pool}}$

Stage 7: Prediction Head

- $\hat{\mathbf{y}} \leftarrow \text{MLP}_{\text{head}}(\mathbf{z})$
- return $\hat{\mathbf{y}}$

Multi-Head Attention

$$\text{MHA}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{Concat}([\text{Attn}(\mathbf{QW}_1^K, \mathbf{KW}_2^K, \mathbf{VW}_3^V)])W_4^V$$

Scaled Dot-Product Attention:

$$\text{Attn}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{QK}^T}{\sqrt{d_{\text{head}}}}\right)\mathbf{V}$$

SwiGLU Feed-Forward:

$$\text{FF}_{\text{SwiGLU}}(\mathbf{h}) = (\text{SiLU}(\mathbf{hW}_1) \odot \mathbf{hW}_2)W_3$$

MLP block (used for gate/feature/head):

$$\text{MLP}_k(\mathbf{x}) = \mathbf{W}_1^{(k)} \text{Dropout}(\text{Gelu}(\mathbf{W}_2^{(k)} \text{LN}(\mathbf{x})))$$

Architecture constants:

$$\text{MLP}_{\text{gate}}: \mathbb{R}^{2d} \rightarrow \mathbb{R}^{2d} \text{ (hidden } d/2), \quad \text{MLP}_{\text{fusion}}: \mathbb{R}^{2d} \rightarrow \mathbb{R}^{2d} \text{ (hidden } d),$$
$$d = 806, \quad H = 8, \quad d_{\text{head}} = 112, \quad d_{\text{hidden}} = \frac{8d}{N}, \quad N = 3$$

Results

Gene Ontology									
(Main Task)	Molecular Function (MF)			Biological Process (BP)			Cellular Component (CC)		
Method	AUPR _{macro}	F _{max}	S _{min}	AUPR _{macro}	F _{max}	S _{min}	AUPR _{macro}	F _{max}	S _{min}
BLAST	0.136	0.328	0.632	0.067	0.336	0.651	0.097	0.448	0.628
FunFams	0.367	0.572	0.531	0.260	0.500	0.579	0.288	0.672	0.503
DeepGO	0.391	0.577	0.472	0.182	0.493	0.577	0.263	0.549	0.550
DeepPFI	0.495	0.625	0.437	0.261	0.540	0.543	0.274	0.613	0.527
PfResGO	0.602	0.692	0.417	0.293	0.568	0.535	0.361	0.674	0.498
HEAL-PDB	0.571	0.691	0.401	0.259	0.565	0.540	0.342	0.655	0.501
GGN-GO-PDB	0.601	0.698	0.400	0.291	0.643	0.546	0.347	0.657	0.510
Omni-Prot (ours)	0.658	0.743	0.342	0.342	0.615	0.498	0.453	0.728	0.420
† BLAST: Transfers protein functions via nearest-neighbor local sequence alignment; commonly used as a classical homology-transfer baseline.									
† FunFams: Subdivides CATH homologous superfamilies into functionally coherent FunFams for domain-level function annotation/transfer.									
† DeepGO: An early deep-learning approach that learns sequence and cross-species protein-protein interaction features while explicitly modeling Gene Ontology hierarchical dependencies.									
† DeepPFI: Combines protein language model sequence embeddings with graph neural networks over residue-residue contact maps to predict GO terms.									
† PFIresGO: An attention-based framework that integrates Gene Ontology graph relationships with representations extracted from protein language models.									
† HEAL: Employs a hierarchical graph transformer trained with contrastive learning on structural graphs, using protein language model embeddings as node features for function prediction.									
† GGN-GO: A geometric graph neural network leveraging protein language model embeddings and multiscale structural representations (atomic and residue levels) through geometric vector perceptors; the prior state-of-the-art, with heavy reliance on structural inputs rather than fully exploiting sequence-based foundation models.									