

Omni-Prot: Unifying Protein Function Prediction with PLM-Centric Fusion

Introduction

Proteins are the molecular workhorses of life, orchestrating virtually every biological process from catalyzing reactions to transmitting signals and maintaining cellular structure. Understanding protein function from amino acid sequence alone remains one of the most fundamental challenges in biology. This challenge has direct implications across medicine, industry, and basic research, yet the vast majority of known proteins remain functionally uncharacterized.

Applications

Understanding protein function has profound implications across virtually every domain of biological research and biotechnology.

- Drug Discovery**
 - Identifying therapeutic targets and predicting off-target effects by understanding which proteins participate in disease-relevant pathways
- Precision Medicine**
 - Interpreting variants of unknown significance by predicting whether mutations disrupt protein function in patient genomes
- Enzyme Engineering**
 - Screening metagenomic databases to identify promising biocatalysts for industrial applications including plastic degradation, carbon capture, and sustainable biofuel
- Functional Genomics**
 - Annotating novel protein sequences from microbiome studies to reveal ecological roles and therapeutic potential hidden in environmental samples

The Dark Proteome

Despite the critical importance of understanding protein function, a massive gap exists between available sequence data and functional knowledge—creating a fundamental bottleneck limiting human progress against disease, environmental challenges, and industrial applications.

- High-throughput sequencing has cataloged ~246 million protein sequences in UniProt as of 2024
- Only ~574,000 entries (<0.25%) have experimentally validated functions through manual curation in Swiss-Prot
- The remaining 99.75% constitute the "dark proteome"—molecular machines whose therapeutic potential, ecological roles, and disease associations remain invisible to science
- The number of uncharacterized proteins continues to grow, adding approximately three hundred new sequences per minute
- Experimental characterization proceeds far too slowly to close this gap: a single protein's function can take weeks to months of laboratory work to determine

Traditional Computational Approaches

Consequently, the field has turned to computational prediction as the only scalable solution to bridge this widening gap between sequence availability and functional understanding.

BLAST: For over three decades, the most widely used tool for protein annotation; identifies sequences similar to already-characterized proteins and transfers their functional labels via homology-based inference

Limitations of homology transfer:

- Functional annotation is typically shallow even when similarity is plausible
- Can only annotate proteins that closely resemble already-characterized sequences
- Leaves the vast majority of the dark proteome—especially proteins from understudied organisms, extreme environments, and metagenomic discoveries—functionally invisible

Current Deep Learning Methods

More recently, deep learning methods have shown promise by learning latent relationships between sequences and functions that transcend simple sequence similarity, yet these approaches face their own critical limitations. Three key limitations persist.

- Prevailing systems compress rich embeddings through aggressive low-dimensional bottlenecks, rely on a single PLM despite complementary inductive biases across architectures, or employ prediction frameworks that fail to fully express PLM representational capacity
- The field frequently overemphasizes structural inputs and places undue weight on auxiliary evolutionary information such as multiple sequence alignments (MSAs)
- Requirements for predicted structures from computationally expensive pipelines (e.g., AlphaFold) or evolutionary alignments from database searches add substantial computational overhead, increase inference latency, limit applicability to novel proteins lacking evolutionary history, and risk propagating errors from upstream predictions

Research Objectives

- Investigate the true drivers of performance gains in protein function prediction through systematic ablation experiments—isolating the contributions of auxiliary inputs (3D coordinates, MSAs) versus protein language model representations
- Design a PLM-centric architecture that fully leverages complementary protein language models through advanced fusion techniques (cross-attention, adaptive gating) and training strategies validated through systematic experimentation
- Evaluate rigorously on Gene Ontology prediction across all three ontologies (MF, BP, CC) against state-of-the-art methods, with additional benchmarking on enzyme classification and stability regression to assess generalization and fine-tuning capability
- Develop an accessible, serverless web application with a user-friendly interface for sequence-to-function prediction, alongside open-source implementation enabling mass annotation at scale

Protein Data Modality

Proteins are linear polymers defined entirely by their amino acid sequence—a chain of residues drawn from a 20-letter alphabet that encodes all the information necessary to determine a protein's structure, function, and behavior. This primary structure is the most fundamental and universally available representation of any protein.

VRDIAKPCNVEYCGITQDCKNLCTENGASGYCQWGGKYGNACWCILKPLDPSVIRPVGKCCQ

In modern deep learning, raw sequences are tokenized and passed through protein language models (PLMs)—large neural networks pretrained on millions/billions of protein sequences via self-supervised objectives such as masked language modeling—to produce rich, per-residue feature vectors (embeddings) that capture latent biochemical and evolutionary properties far beyond what raw one-hot encodings provide.

Auxiliary Information

3D Structure (Coordinates)

- Experimentally determined (e.g., Protein Data Bank) or computationally predicted (e.g., AlphaFold) three-dimensional atomic coordinates, represented as Ca backbone positions
- Often encoded as residue contact graphs or geometric features, providing explicit spatial proximity and geometric constraint information

Multiple Sequence Alignments (MSAs)

- Collections of evolutionarily related sequences aligned position-by-position, generated by searching large sequence databases for homologs
- Encode co-evolutionary information: correlated mutations across aligned positions reveal residue-residue contacts and functional constraints shaped by natural selection

Additional Modalities

- Other modalities include InterPro domain annotations, co-evolutionary coupling matrices, protein-protein interaction networks, and secondary structure
- These are not explicitly ablated in this work, as 3D coordinates and MSAs subsume much of the signal these sources provide, making them the most informative test cases for evaluating auxiliary feature dependence

Datasets & Evaluation Metrics

Gene Ontology

Task Background

- The Gene Ontology (GO) provides a controlled vocabulary for describing protein function across three orthogonal ontologies: molecular function (MF), biological process (BP), and cellular component (CC)
- MF captures specific biochemical activities (e.g., ATP binding, oxidoreductase activity); BP describes broader pathways (e.g., DNA repair, signal transduction); CC localizes proteins to subcellular structures (e.g., cytosol, ribosome)
- Technically challenging: labels are hierarchical, exhibit long-tailed frequency distributions, and require models to generalize to remote homologs and rare terms while remaining ontology-consistent
- Central task targeted in this work—regarded as one of the most important benchmarks in protein informatics since it is well-defined while addressing the key overall aspects of protein function through its three ontologies

Dataset

- PDBCh: 36,641 protein chains from the Protein Data Bank clustered at 95% sequence identity, split 8:1:1 into 29,313 training, 3,664 validation, and 3,664 test proteins
- Label space: 489 MF terms, 1,943 BP terms, 320 CC terms—exhibiting moderate label imbalance typical of real-world annotation databases (rare terms ~50 examples, common terms in thousands)
- Test set remains fixed across all three ontologies for standard cross-study comparability

Evaluation Metrics

- GO predictions undergo standard hierarchy-aware propagation preprocessing to ensure consistency with the GO ontology tree
- Protein-centric F_{max} : Best thresholded trade-off between per-protein precision and recall; higher means balanced, reliable hits at the protein level
- Function-centric $AUPR_{macro}$: Averages average-precision over GO terms so rare and common labels count equally; higher means better performance on long-tailed terms
- Semantic distance S_{dist} : Minimizes information-weighted error on the GO DAG; lower values indicate more semantically accurate and hierarchy-consistent predictions

$$F_{max} = \max_{\tau \in \mathcal{T}} \frac{2\bar{P}(\tau)\bar{R}(\tau)}{\bar{P}(\tau) + \bar{R}(\tau)}$$
$$AUPR_{macro} = \frac{1}{|\mathcal{L}|} \sum_{t \in \mathcal{L}} AP_t$$
$$S_{dist} = \min \sqrt{RU(r)^2 + MI(r)^2}$$

Enzyme Commission

Task Background

- EC number: hierarchical four-level taxonomy specifying the chemical transformation catalyzed by an enzyme (class → subclass → sub-subclass → substrate specificity)
- Near single-label classification with extremely long-tailed frequency distribution

Dataset

- Training: 74,487 sequences clustered at 70% identity from UniProt; Validation: NEW-392 (392 Swiss-Prot enzymes); Test: Price-149 (149 experimentally validated enzymes)

Evaluation Metric

- AUC: Area under ROC curve measuring class separability; higher means stronger discrimination
$$AUC = \int_0^1 TPR(FPR) d(FPR), \quad TPR = \frac{TP}{TP + FN}, \quad FPR = \frac{FP}{FP + TN}$$

Stability

Task Background

- Protein stability: ability of a sequence to retain its native fold under stress (heat, denaturants, proteolysis), summarized as a continuous score
- Regression task probing mutational landscape; models must generalize beyond single-mutation trends

Dataset

- TAP stability benchmark: 53,416 training, 2,512 validation, 12,851 test sequences from a high-throughput do-no-misprotein design campaign

Evaluation Metric

- Spearman's ρ : Rank correlation between predicted and true stability scores; higher means stronger monotonic agreement
$$\rho = \frac{\text{Cov}(r(s), r(y))}{\sigma_{r(s)} \sigma_{r(y)}} = 1 - \frac{6 \sum_{i=1}^n (r(s_i) - r(y_i))^2}{n(n^2 - 1)}$$

Input Modality Investigation

Multiple Sequence Alignments

Background & Motivation

- Multiple sequence alignments (MSAs) capture evolutionary relationships by aligning homologous sequences, revealing conservation patterns and co-evolutionary signals
- Key question: Do MSAs provide measurable benefit to protein function prediction tasks when strong protein language models are available?

Experimental Setup

- Compared two architecturally identical models on MF prediction with varied inputs:
 - Baseline: Four PLM embeddings (ESM-C, ProtT5, Ankh3-XL, PGLM)
 - MSA-augmented: Pooled MSA features replace PGLM as fourth stream
- Substitution design ensures identical capacity, isolating effect of auxiliary evolutionary information vs single-sequence modeling
- MSAs encoded via MSA Transformer with learned conservation-weighted pooling (see diagram):

$$x_{seq} = W_{MLM} \cdot M_{seq}, \quad \alpha_{seq} = \frac{\exp(\frac{1}{\sqrt{d}} \cdot \frac{1}{N})}{\sum_{i=1}^N \exp(\frac{1}{\sqrt{d}} \cdot \frac{1}{N})}, \quad h_i = \sum_{j=1}^N \alpha_{seq} M_{seq}^{(j)}$$

Structural Information

A. KNN GRAPH CONSTRUCTION

Real Coordinates → Random Coordinates

B. EGNN LAYER (x4)

1. Distance Features

$$d_{ij} = |q_i - q_j|$$

C. READOUT PIPELINE

Processed Node Features → Masked Mean Pooling → Protein Embedding → Classification Head

Masked Mean Pooling: $g = \sum_i h_i / N$

Protein Embedding: $g \in \mathbb{R}^{128}$

Classification Head: MF/BP/CC predictions

D. Ablation Results

Graph Ablation: MF AUPR_{macro} over Epochs

Graph Ablation: BP AUPR_{macro} over Epochs

Graph Ablation: CC AUPR_{macro} over Epochs

Scaling Protein Language Models

Background & Motivation

- Nearly all existing prediction methods rely on a single PLM to extract sequence features
- Key question: Does incorporating multiple PLMs with diverse architectural biases improve performance?

Epoch	1	2	3	4	6
1	0.146	0.272	0.287	0.341	0.355
2	0.389	0.540	0.601	0.646	0.644
3	0.553	0.675	0.744	0.760	0.777
4	0.632	0.745	0.804	0.809	0.822
5	0.710	0.795	0.842	0.845	0.845
6	0.754	0.815	0.856	0.861	0.861
7	0.791	0.837	0.867	0.877	0.875

Architecture

Diagram

Input: Protein Sequence

Stage 1: Embedding

Stage 2: Pre-trained Language Models (ESM-C, Ankh3-XL, ProtT5, PGLM)

Stage 3: Intra-stream fusion

Stage 4: Inter-stream cross-attention

Stage 5: Bidirectional Cross-Attention

Stage 6: Masked Mean Pooling

Stage 7: MLP Head

Stage 8: Gene Ontology Terms

Stage 9: Predictions

Stage 10: Additional Tasks

Results

Gene Ontology

Method	Molecular Function (MF)			Biological Process (BP)			Cellular Component (CC)		
	AUPR _{macro}	F _{max}	S _{dist}	AUPR _{macro}	F _{max}	S _{dist}	AUPR _{macro}	F _{max}	S _{dist}
BLAST	0.136	0.328	0.632	0.067	0.336	0.651	0.097	0.448	0.628
FunFams	0.367	0.572	0.531	0.260	0.500	0.579	0.288	0.672	0.503
DeepGO	0.391	0.577	0.472	0.182	0.493	0.577	0.263	0.649	0.550
DeepFRI	0.495	0.625	0.437	0.261	0.540	0.543	0.274	0.613	0.527
PFresGO	0.602	0.692	0.417	0.293	0.568	0.535	0.361	0.674	0.498
HEAL-PDB	0.571	0.691	0.401	0.259	0.565	0.540	0.342	0.655	0.501
GGN-GO-PDB	0.601	0.698	0.400	0.291	0.643	0.546	0.347	0.657	0.510
Omni-Prot (ours)	0.658	0.743	0.342	0.342	0.615	0.498	0.453	0.728	0.420

Enzyme Commission

Method	AUC
GRAPPI	0.5095
ECPICK	0.5888
CLEAN	0.7334
GraphEC	0.8404
Omni-Prot (ours)	0.855

Stability

Method	Spearman's ρ	Relative to the strongest prior, ProteinBERT (0.76), Omni-Prot reaches a ρ of 0.85 ($\Delta = +0.09$; +11.8%). Gains on other widely used baselines are consistent: Transformer (0.73) $\rightarrow \Delta = +0.12$; +16.4%; LSTM (0.69) $\rightarrow \Delta = +0.16$; +23.2%; RTLS (0.71) $\rightarrow \Delta = +0.14$; +19.7%; k-NN (0.21) $\rightarrow \Delta = +0.64$; +305%, and One-Hot (0.19) $\rightarrow \Delta = +0.66$; +347%.
One-Hot	0.19	
k-NN	0.21	
LSTM	0.69	
RT (L2C)	0.71	
Transformer	0.73	
UniProt	0.76	
ProteinBERT	0.85	
Omni-Prot (ours)	0.85	

Supplemental Task

Method	Spearman's ρ	Relative to the strongest prior, ProteinBERT (0.76), Omni-Prot reaches a ρ of 0.85 ($\Delta = +0.09$; +11.8%). Gains on other widely used baselines are consistent: Transformer (0.73) $\rightarrow \Delta = +0.12$; +16.4%; LSTM (0.69) $\rightarrow \Delta = +0.16$; +23.2%; RTLS (0.71) $\rightarrow \Delta = +0.14$; +19.7%; k-NN (0.21) $\rightarrow \Delta = +0.64$; +305%, and One-Hot (0.19) $\rightarrow \Delta = +0.66$; +347%.
One-Hot	0.19	
k-NN	0.21	
LSTM	0.69	
RT (L2C)	0.71	
Transformer	0.73	
UniProt	0.76	
ProteinBERT	0.85	
Omni-Prot (ours)	0.85	

State-of-the-art across MF and CC on all metrics; BP shows strong ranking improvement (AUPR_{macro}) with competitive F_{max}

Molecular Function: AUPR_{macro} +9.35% (0.602→0.658), F_{max} +6.50% (0.698→0.743), S_{dist} -14.4% (0.400→0.342)—gains in MF directly translate to more accurate identification of specific enzymatic activities and binding capabilities, beneficial for drug target identification

Biological Process: AUPR_{macro} +16.7% (0.293→0.342), S_{dist} -7.0% (0.535→0.498)—substantial AUPR_{macro} improvement indicates markedly better ranking of rare biological processes, suggesting the model captures complex pathway-level relationships more effectively despite BP's inherent difficulty in threshold-dependent metrics

Cellular Component: AUPR_{macro} +25.4% (0.200→0.453), F_{max} +8.05% (0.674→0.728), S_{dist} -15.7% (0.498→0.420)—dramatic AUPR_{macro} gains reflect substantially improved discrimination of rare subcellular compartments and protein complexes, particularly valuable for understanding cellular organization in understudied organisms

Consistent AUPR_{macro} improvements across all ontologies demonstrate superior handling of long-tailed label distributions—critical for real-world annotation where rare terms are often most biologically interesting

Achieved using sequence-only input—outperforms methods relying on structural geometry or evolutionary alignments, validating PLM-centric design

Enzyme Commission

Method	AUC
GRAPPI	0.5095
ECPICK	0.5888
CLEAN	0.7334
GraphEC	0.8404
Omni-Prot (ours)	0.855

Against the strongest prior, GraphEC (0.8404), Omni-Prot attains an AUC score of 0.855 ($\Delta = +0.0146$; +1.74%). Gains are larger versus earlier baselines: CLEAN (0.7334) $\rightarrow \Delta = +0.1216$; +16.58%, ECPICK (0.5888) $\rightarrow \Delta = +0.2662$; +45.21%, and GRAPPI (0.5095) $\rightarrow \Delta = +0.3455$; +67.81%.

Prior Methods

- GRAPPI: Constructs a weighted protein graph from InterPro domain-composition similarity and propagates EC labels from neighboring proteins
- ECPICK: Uses evidential deep learning to generate EC predictions with calibrated uncertainty, improving reliability and trustworthiness on large-scale enzyme databases
- CLEAN: Applies contrastive learning to protein language model representations using curated positive and negative pairs, enabling accurate EC annotation and identification of mislabeled or promiscuous enzymes
- GraphEC: A geometric graph-based framework that combines ESMFold-predicted structures with protein language model embeddings, localizes active sites, and refines EC predictions through label diffusion

Experimental Setup

- Trained (N)-equivariant graph neural network (EGNN) on all three gene ontologies with described varied inputs, comparing real and random coordinates (see diagram)
- Nodes initialized with ESM-C PLM embeddings to purely isolate the impact of the geometry itself

Results & Interpretation

- MF: 0.696 real vs 0.705 random | BP: 0.367 real vs 0.374 random | CC: 0.553 real vs 0.565 random
- Random coordinates achieve nearly identical performance across all ontologies
- Interpretation: PLMs have internalized high-level spatial proximity patterns through co-occurrence statistics for protein function prediction—explicit 3D geometry provides negligible additional signal

Web Tool & Accessibility

Omni-Prot Parameter Overview

Input Sequence:

Model:

ESM-C:

Ankh3-XL:

ProtT5:

PGLM:

Stream 1 Encoder:

Stream 2 Encoder:

Masked Mean Pooling:

MLP Head:

Gene Ontology Terms:

Predictions:

Additional Tasks:

Frontend

- React + TypeScript
- User interface
- Input validation
- Results display
- Deployment on Vercel

Backend API

- Model Services
- Python functions
- Request handling
- Results display
- Model inference
- Auto-scaling

ML Model

- Omni-Prot
- Protein language models
- GO term prediction
- Function analysis

Results

- GO Predictions
- GO terms
- Confidence scores
- MF/BP/CC functions
- Visual display

Key Advantages

- Released user-friendly web application for immediate practical use, designed primarily for GO term prediction where the approach shows strongest performance improvements
- Many existing deep learning models for protein function prediction lack accessible web interfaces for practical deployment, limiting use by broader research community
- Current web-based protein function tools rely on BLAST and homology transfer, which cannot learn functional patterns beyond sequence similarity to close database matches
- Model overcomes these limitations while providing massive improvements compared to existing homology-based methods
- Particularly useful for annotating metagenomic discoveries, interpreting disease variants, and accelerating pathway design (within supported GO dictionary)
- Tool optionally supports EC classification and stability prediction for enzyme engineering applications
- Enables scalable mass annotation of proteins for large-scale analyses due to model's sole reliance on amino acid sequences—a critical capability for high-throughput proteomics and functional genomics studies
- Modular architecture enables researchers to adapt framework to new tasks through fine-tuning on custom datasets
- Researchers can seamlessly integrate more powerful PLMs as they emerge, ensuring long-term relevance without architectural redesign

Conclusion

Future Impact

- Drug discovery: GO predictions for molecular function and biological process directly inform target prioritization—identifying which uncharacterized proteins participate in disease-relevant pathways while flagging potential off-target effects through predicted functional similarities
- Precision medicine: Cellular component and molecular function predictions provide interpretive context for variants of unknown significance—enabling clinicians to assess whether mutations likely disrupt critical protein functions or localization
- Enzyme engineering: Sequence-only inference enables rapid screening of massive metagenomic databases for promising biocatalysts when evaluating candidates for plastic degradation, carbon capture, or sustainable biofuel applications
- Metagenomics: Massive improvements over BLAST enable functional annotation of novel sequences from understudied organisms and extreme environments—proteins that homology-based methods leave functionally invisible
- Scalable deployment: Independence from MSA generation and structure prediction enables true high-throughput analysis—annotating thousands of proteins simultaneously without computational bottlenecks
- Extensibility: Modular architecture allows seamless integration of more powerful PLMs as they emerge—ensuring long-term relevance without architectural redesign or retraining from scratch

Limitations

- Computational constraints prevented extensive repeated trials across all experimental configurations—though preliminary experiments showed low variance
- EC dataset exhibits label inconsistencies between training and test splits and extreme label imbalance—may affect absolute performance metrics for the secondary task
- Performance on proteins from deeply understudied taxonomic lineages, extreme environments, or synthetic designs diverging substantially from natural evolutionary trajectories may vary
- Future work should systematically evaluate robustness across underrepresented protein spaces

Conclusion

- Omni-Prot achieves state-of-the-art on GO prediction (all three ontologies), EC classification, and stability regression using sequence-only input
- Auxiliary inputs (3D coordinates, MSAs) provide negligible benefit when strong PLM representations are properly leveraged
- Future research should focus on extracting information and increasing capture performance of protein language models to maximize performance on a wide array of protein function prediction tasks
- Multi-PLM fusion with complementary architectural biases drives performance—validated through systematic ablation
- Sequence-only design eliminates error propagation from upstream pipelines, simplifies deployment, and enables greater potential robustness in terms of generalization to proteins lacking evolutionary history
- Open-source implementation and accessible web interface make advances immediately available to the broader research community